



CONTACT

Madrid, Spain
rosalinda.barral@gmail.com
+34 618 626 596
linkedin.com/in/rosalinda-barral
github.com/rosalinda75

SKILLS

Python PyTorch
scikit-learn TensorFlow
XGBoost MLflow
Apache Airflow
Google Cloud Platform
Vertex AI BigQuery SQL
Docker Kubernetes Git
GitHub Actions FastAPI
Pandas NumPy Feast
Tableau

LANGUAGES

Spanish – Native
English – C1
French – B1

CERTIFICATIONS

Google Cloud Professional Machine Learning Engineer
Google Cloud Professional Data Engineer

Rosalinda Barral

Senior Machine Learning Engineer

PROFILE

Machine Learning engineer with eight years of end-to-end experience designing and deploying predictive systems in banking, travel, and consulting environments. Expertise in developing production-grade ML pipelines on GCP using PyTorch, TensorFlow, MLflow, and Airflow, delivering measurable impact on revenue, latency, and operational cost. Proven ability to lead cross-functional teams, mentor junior engineers, and translate business requirements into scalable AI solutions.

EXPERIENCE

Senior Machine Learning Engineer, Capgemini Invent

09/2023 – Present

Madrid, Spain

- Led migration of client ML workloads to Vertex AI, cutting average model training time from 4 hours to 25 minutes via distributed training and reducing GPU costs by 35%.
- Designed a centralized MLflow tracking server that now manages 50+ model versions, decreasing model promotion lead time from 3 days to 4 hours.
- Optimized inference pipelines with TensorRT and GPUs, achieving a 40% latency reduction for real-time recommendation APIs serving 2 M requests/day.
- Implemented a feature store using Feast on GCP, enabling consistent feature serving across training and online environments for 6 data science teams.
- Mentored a squad of 4 ML engineers and 2 data analysts, improving team velocity by 22% as measured by sprint story points delivered.

Machine Learning Engineer, eDreams ODIGEO

09/2020 – 08/2023

Madrid, Spain

- Built a hotel demand forecasting model using Prophet and XGBoost that improved forecast MAPE from 12% to 7%, driving a 3.5% revenue uplift (~€4.5M annually).
- Created Airflow DAGs orchestrating nightly ETL of 2 TB of clickstream data into BigQuery, cutting data latency from 6 h to 45 min.
- Deployed real-time pricing models on Vertex AI endpoints, handling peak loads of 8 k QPS with 99.9% uptime.
- Introduced MLflow experiment tracking, reducing model iteration cycles from 2 weeks to 3 days and increasing the number of tested variants per sprint by 3x.
- Collaborated with a team of 5 data engineers and 3 product managers to launch a personalized email campaign engine, increasing conversion rate by 1.8 points.

Machine Learning Engineer, Banco Santander

09/2017 – 08/2020

Madrid, Spain

- Developed a PyTorch-based fraud detection model that reduced false positives by 18% and saved an estimated €2.3 M per year in avoided losses.
- Containerized model serving with Docker and Kubernetes on GKE, achieving a peak throughput of 15 k transactions per second with <30 ms latency.
- Built SQL-based feature extraction pipelines in BigQuery, processing 500 M daily transaction records and feeding real-time scoring services.

- Integrated MLflow for model versioning, cutting model rollback time from hours to minutes during production incidents.
- Worked within a squad of 3 ML engineers, 2 data analysts, and 4 risk specialists to deploy the model across three European markets.

EDUCATION

Universidad Complutense de Madrid

BSc, Mathematics · 2013–2017

Universidad Politécnica de Madrid

MSc, Data Science · 2017–2019