



CONTACT

Bengaluru, India
yuvraj.parsa@gmail.com
+91 98778 76051
linkedin.com/in/yuvraj-parsa
github.com/yuvraj62

SKILLS

Apache Spark · Scala · Python ·
Apache Airflow · AWS (EC2, S3,
Glue, Redshift, MWAA) · Snowflake
· dbt · Apache Kafka · Docker ·
Kubernetes · Git · SQL · Tableau ·
Great Expectations · Linux

LANGUAGES

English – C2
Hindi – Native
Kannada – Native

CERTIFICATIONS

AWS Certified Data Analytics –
Specialty (2023)
Google Cloud Professional Data
Engineer (2024)
Snowflake SnowPro Core
Certification (2022)

Yuvraj Parsa

Data Engineer | 5 years experience in Spark, Scala, and AWS data platforms

PROFILE

Data engineer with five years of end-to-end experience building scalable data pipelines on AWS using Apache Spark, Scala, Python, Airflow, and Snowflake. Proven track record of reducing latency, cutting costs, and enabling data-driven decision making for fintech and enterprise clients. Adept at mentoring teams and implementing data quality frameworks.

EXPERIENCE

Senior Data Engineer, PhonePe 2022 – Present
Bengaluru, India

- Designed and executed a Spark Structured Streaming pipeline ingesting 2.5 TB/day of clickstream data from Kafka into Snowflake, reducing end-to-end latency from 45 minutes to under 5 minutes.
- Migrated 30+ legacy ETL jobs from on-prem Airflow to AWS Managed Workflows for Apache Airflow (MWAA), cutting operational overhead by 40% and saving approximately ₹1.2 million annually.
- Built dbt models that transformed raw transactional data into a unified customer 360 view, enabling PhonePe’s marketing team to increase campaign ROI by 18% year-over-year.
- Implemented schema evolution checks using Great Expectations, preventing 98% of data quality incidents in production.
- Mentored a team of four junior engineers, conducting bi-weekly code reviews and reducing average PR turnaround time from three days to twelve hours.

Data Engineer, Tata Consultancy Services (TCS) 2020 – 2022
Bengaluru, India

- Developed Scala-based Spark batch jobs processing 1.2 TB of daily transaction logs for a banking client, improving job runtime from 90 minutes to 30 minutes via partition pruning and broadcast joins.
- Created Python Airflow DAGs orchestrating data extraction from Oracle, transformation in Spark, and loading into AWS Redshift, achieving 99.5% SLA compliance over 12 months.
- Integrated Kafka Connect connectors to stream CDC updates from MySQL to S3, enabling near-real-time analytics dashboards used by 150+ business users.
- Optimized Snowflake warehouse sizing using query profiling, reducing monthly compute spend by 22% (≈₹850K).
- Authored internal documentation and conducted workshops on Scala best practices for a squad of six engineers, increasing code reuse by 35%.

EDUCATION

National Institute of Technology Tiruchirappalli
BTech, Computer Science · 2016–2020